

WHITE PAPER

# AIRI CONFIGURATION GUIDE

SCALE-OUT AI-READY INFRASTRUCTURE  
ARCHITECTED BY PURE STORAGE AND NVIDIA

# TABLE OF CONTENTS

<b>INTRODUCTION</b>	3
<b>HOW TO USE THIS GUIDE</b>	4
<b>SYSTEM ARCHITECTURE</b>	4
<b>PREREQUISITES</b>	5
Install Additional Package Dependencies	5
<b>NETWORK CONFIGURATION</b>	5
Switch Configuration	6
DGX-1 Configuration	6
Configure Network Interfaces	7
TCP Traffic Evaluation	9
RDMA Traffic Evaluation	9
FlashBlade Configuration	11
<b>ADDITIONAL RESOURCES</b>	12
<b>APPENDIX A: HOROVOD CONTAINER IMAGE</b>	12

## INTRODUCTION

Recent advances in AI and deep learning hold tremendous promise for a new wave of innovation for enterprises, turning their data into intelligent applications and products. Researchers have fueled these advancements using high-performance GPUs and leveraging massive growth in available datasets. While enterprises seek to get started on their AI journey, they are stuck with legacy technologies like CPUs and spinning disk, and the complexities of building an AI-ready infrastructure with them.

Designing, configuring, and maintaining infrastructure to satisfy the challenges of large-scale deep learning requires a significant investment of time and resources to avoid unforeseen delays, bottlenecks, or downtime. Engineers at NVIDIA® and Pure Storage® collaborated to address many of the complexities that come with scale-out infrastructure for AI workloads.

AIRI™ is the industry's first complete AI-ready infrastructure, architected by Pure Storage and NVIDIA to extend the power of NVIDIA® DGX™ systems, enabling AI-at-scale for every enterprise. AIRI is a converged infrastructure stack, purpose-built for large-scale deep learning environments. The entire stack is configured and tested as a complete solution, avoiding the intricate configuration and tuning required otherwise.

To learn more, please refer to the AIRI Reference Architecture at <https://github.com/PureStorage-OpenConnect/AIRI>.




FIGURE 1. AIRI

## HOW TO USE THIS GUIDE

This guide describes how to configure an AIRI system containing 4x NVIDIA DGX-1 servers and a Pure Storage FlashBlade. To optimize DGX-1 node-to-node communication, the configuration uses an RDMA over Converged Ethernet (RoCE) fabric. The same high-performance fabric carries storage traffic from FlashBlade to the DGX-1 servers, simplifying system configuration and deployment.

An Installation Engineer is available to assist with installing and configuring the system and can help with questions or assistance during setup.

## SYSTEM ARCHITECTURE

The AIRI architecture is depicted in the diagram below, along with a description of the core components of the system.

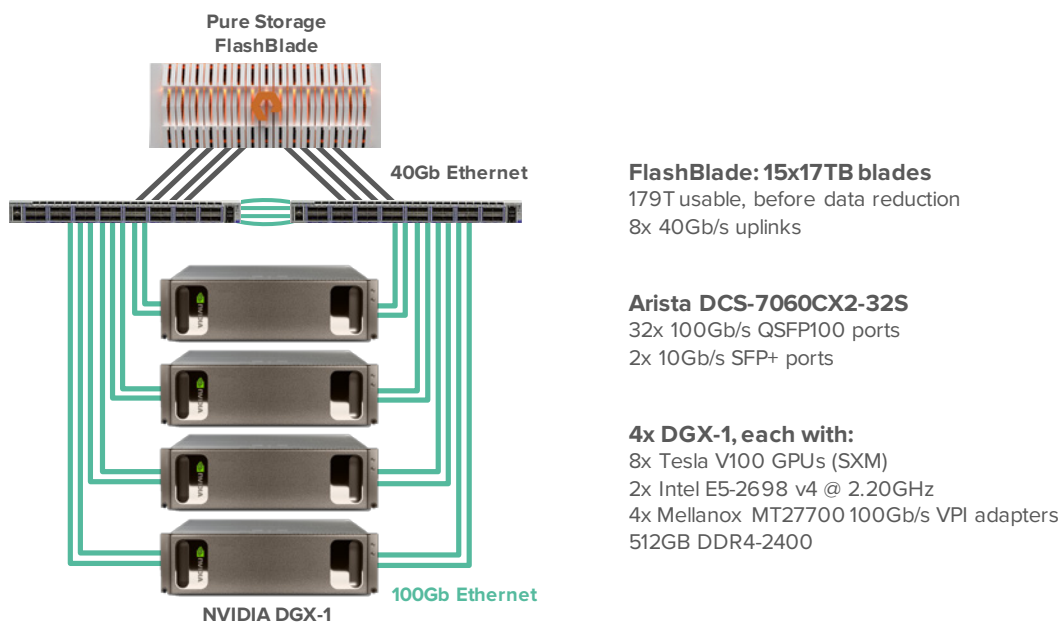


FIGURE 2. AIRI architecture

The AIRI system is designed for scale-out deep learning workloads, and is not restricted to this ratio of storage to compute. As datasets and workload requirements scale, additional DGX-1 servers can be provisioned and instantly access all available data. Similarly, as storage capacity or performance demands grow, additional blades can be added to the FlashBlade system with zero downtime or re-configuration. Commands in this guide detail the configuration of 4x DGX-1 servers but can be easily modified to configure a larger AIRI system. Please reach out to Pure Storage when planning deployments with > 10 DGX-1 servers per FlashBlade array.

## PREREQUISITES

Configure DGX-1 servers and FlashBlade according to their respective setup guides. For the initial setup, configure the 10Gb/s connections on the DGX-1 servers for management, and create a subnet and VIP on FlashBlade for this network.

For the remainder of this guide, we assume four DGX-1 servers have been configured and given names (via /etc/hosts) dgx-X, where X is 1-4.

## Install Additional Package Dependencies

These packages are necessary for the subsequent configuration of vlan networking and TCP/IP network testing.

```
$ sudo apt-get update
$ sudo apt-get install -y vlan iperf3
```

## NETWORK CONFIGURATION

The diagram below shows the overall network topology, excluding a separate management network. In the diagram, FM-1/2 are the two internal fabric modules of aFlashBlade, SW-1/2 are dual 100G switches, and DGX-A, DGX-B, DGX-C, and DGX-D are the DGX-1 servers. Each of the numbers on SW-1/2 indicate the Ethernet port number, and the numbers on the DGX indicate which of the ethernet devices (enp5s0, enp12s0, etc.).

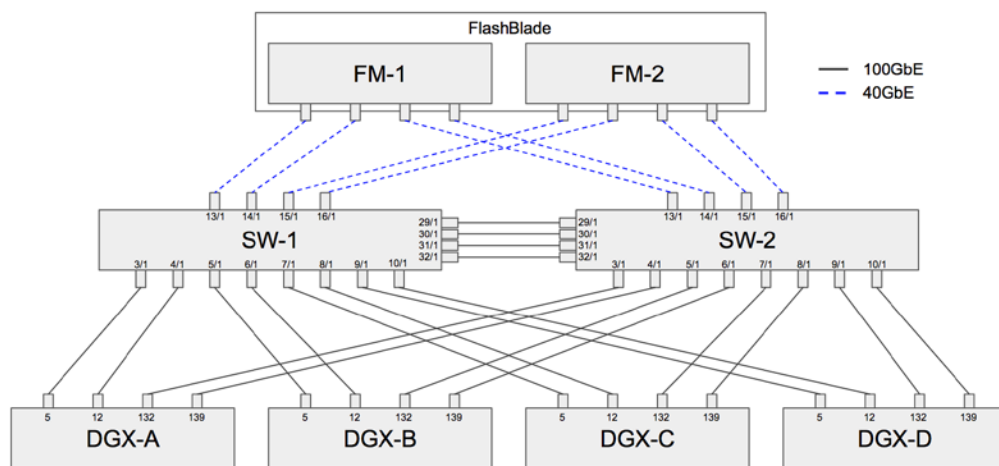


FIGURE 3. Overall network topology

Our reference configuration used Arista 100G, 32-port Ethernet switches (model DCS-7060CX2-32S, running EOS-4.18.2-REV2-FX), but other switch models may be used in the future.

Each DGX is connected into the overall topology by four distinct 100Gb/s links, with two ports connected to each of the two switches. The FlashBlade fabric modules have all 8 of their QSFP uplinks connected, with each fabric module

connecting two ports into each switch. The two switches SW-1 and SW-2 have 4x 100Gb/s connections between them in an MLAG for inter-switch traffic.

There are two logical networks configured on this fabric. First, the DGX-1s use an untagged VLAN (id 1) to carry RoCE traffic. All four 100G ports on each DGX-1 are configured on this VLAN and participate in node-to-node RoCE communications.

Second, there is a TCP network that runs in VLAN 3000. On each DGX-1, two of the 100G ports – enp5s0 and enp132s0 – are selected to carry VLAN 3000. These are placed in an active-backup bond in the Linux bonding driver.

The RoCE traffic is configured in a priority-flow-control class on both the server endpoints and switches, allowing the more latency-sensitive traffic between DGX-1 servers to take priority over the TCP traffic used for storage access.

## Switch Configuration

For each of the 100Gb/s data switches (DCS-7060CX-32S), log into the **switch console** and issue the following commands:

```
SW-1>en
SW-1#config
SW-1(config)#vlan 3000
SW-1(config-vlan-3000)#name tcp
SW-1(config-vlan-3000)#exit
SW-1(config)#int et3/1-et10/1
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#flowcontrol send off
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#flowcontrol receive off
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#switchport trunk allowed vlan 1,3000
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#switchport mode trunk
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#priority-flow-control mode on
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#priority-flow-control priority 3 no-drop
SW-1(config-if-Et3/1,4/1,5/1,...,9/1,10/1)#exit
SW-1(config)#exit
SW-1#exit
```

In addition, add **vlan 3000** to the switchport trunk allowed list for the port-channel connected to FlashBlade.

## DGX-1 Configuration

To switch VPI cards from Infiniband to Ethernet mode, first start the Mellanox Software Tools (MST) set:

```
$ mst start
Starting MST (Mellanox Software Tools) driver set
Loading MST PCI module - Success
```

```
Loading MST PCI configuration module - Success
Create devices
Unloading MST PCI module (unused) - Success
```

Then, change the port type to **Ethernet** (LINK\_TYPE = 2):

```
$ mlxconfig -d /dev/mst/mt4115_pciconf0 set LINK_TYPE_P1=2
Device #1:
-----

Device type:    ConnectX4
PCI device:    /dev/mst/mt4115_pciconf0

Configurations:      Current New
      LINK_TYPE_P1    1      2

Apply new Configuration? ? (y/n) [n] : y
Applying... Done!
-I- Please reboot machine to load new configurations.
```

Repeat for the other three devices:

```
$ mlxconfig -d /dev/mst/mt4115_pciconf1 set LINK_TYPE_P1=2
$ mlxconfig -d /dev/mst/mt4115_pciconf2 set LINK_TYPE_P1=2
$ mlxconfig -d /dev/mst/mt4115_pciconf3 set LINK_TYPE_P1=2
```

Reboot the server to complete. After rebooting, run:

```
$ ibv_devinfo
```

And confirm that all four devices show link\_layer: **Ethernet**.

## Configure Network Interfaces

In the following, we use **iproute2** commands (e.g., `ip link add...`). These are fast to execute and easy to debug, but are not persistent across reboots. Create a persistent version of the network setup by either writing a script to execute the following commands on startup, or create the equivalent using **/etc/network/interfaces** (man 5 interfaces).

The commands below will configure:

- VLAN 3000 virtual interfaces on NICs enp5s0 and enp132s0
- A bond0 active/backup pair over the VLAN3000 interfaces

- A private 198.18.0.0/24 address for bond0
- MTU 9000 (jumbo frames) on all interfaces

Note that the example below uses 198.18.0.11 for one server. Substitute unique addresses for each server in the cluster. All other settings are common across servers. All commands are assumed to be run as root.

```
$ modprobe bonding mode=active-backup miimon=100
$ ip link add name enp5s0.3000 link enp5s0 type vlan id 3000
$ ip link add name enp132s0.3000 link enp132s0 type vlan id 3000
$ ifenslave bond0 enp5s0.3000 enp132s0.3000
$ ip address add 198.18.0.11/24 dev bond0 * unique address / DGX
$ ip link set dev enp5s0 mtu 9000
$ ip link set dev enp12s0 mtu 9000
$ ip link set dev enp132s0 mtu 9000
$ ip link set dev enp139s0 mtu 9000
$ ip link set dev enp5s0.3000 mtu 9000
$ ip link set dev enp132s0.3000 mtu 9000
$ ip link set dev enp5s0 up
$ ip link set dev enp12s0 up
$ ip link set dev enp132s0 up
$ ip link set dev enp139s0 up
$ ip link set dev enp5s0.3000 up
$ ip link set dev enp139s0.3000 up
$ ip link set dev bond0 up
```

Configure the VLAN interfaces used for TCP traffic as **priority 0** (low) egress traffic:

```
$ for I in {0..7}; do vconfig set_egress_map enp5s0.3000 $I 0; done
$ for I in {0..7}; do vconfig set_egress_map enp132s0.3000 $I 0; done
```

And configure the raw VPI devices to use PFC **priority 3**:

```
$ mlx_qos -i enp5s0 --pfc 0,0,0,1,0,0,0,0
$ tc_wrap.py -i enp5s0 -u 3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3
$ mlx_qos -i enp12s0 --pfc 0,0,0,1,0,0,0,0
$ tc_wrap.py -i enp12s0 -u 3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3
$ mlx_qos -i enp132s0 --pfc 0,0,0,1,0,0,0,0
$ tc_wrap.py -i enp132s0 -u 3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3
$ mlx_qos -i enp139s0 --pfc 0,0,0,1,0,0,0,0
$ tc_wrap.py -i enp139s0 -u 3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3
```



## TCP Traffic Evaluation

We use **iperf3** to exercise the 100Gb TCP VLAN (3000) and ensure all DGX-DGX pairs are operating at full network bandwidth for TCP traffic. Because iperf3 is a single-threaded process, we must use multiple client-server pairs to saturate the 100Gb link.

From one of the DGX-1 systems, run two server processes as **daemons**. These commands start two listener processes on ports 5201 and 5202. For the example below, assume we run these on the DGX-1 with IP address 198.18.0.11.

```
$ iperf3 -s -p 5201 -D
$ iperf3 -s -p 5202 -D
```

On another DGX-1, run the following commands concurrently:

```
$ iperf3 -c 198.18.0.11 -t 30 -l 1M -p 5201 -R
$ iperf3 -c 198.18.0.11 -t 30 -l 1M -p 5202 -R
```

These client processes will report total bandwidth sent, similar to the following:

[ ID]	Interval		Transfer	Bandwidth	Retr	
[ 4]	0.00-30.00	sec	161 GBytes	46.2 Gbits/sec	96	sender
[ 4]	0.00-30.00	sec	161 GBytes	46.2 Gbits/sec		receiver

[ ID]	Interval		Transfer	Bandwidth	Retr	
[ 4]	0.00-30.00	sec	174 GBytes	49.9 Gbits/sec	15	sender
[ 4]	0.00-30.00	sec	174 GBytes	49.9 Gbits/sec		receiver

As shown, the total bandwidth sent is 46.2 + 49.9 = 96.1 Gb/s. Repeating these results with all pairs, we confirm the TCP network is operational and able to saturate the bond as configured.

## RDMA Traffic Evaluation

The DGX-1 ships with the OpenFabrics' **perftests** suite for evaluating native IB verbs performance. These tools run exclusively on RDMA-capable networks, supporting both Infiniband and RDMA-over-Converged-Ethernet physical layers.

On one of the DGX-1 systems – in the example below, this is “dgx-1” – start a server process with to run a **bandwidth test**:

```
$ ib_send_bw --report_gbits -aF -d mlx5_0
```

On a second DGX-1 system, run the **client**:

```
$ ib_send_bw --report_gbits -aF -d mlx5_0 dgx-1
```

---

```
                Send BW Test

Dual-port       : OFF           Device       : mlx5_0
Number of qps   : 1             Transport type : IB
Connection type : RC            Using SRQ      : OFF
TX depth        : 128
CQ Moderation   : 100
Mtu             : 4096[B]
Link type       : Ethernet
GID index       : 3
Max inline data : 0[B]
rdma_cm QPs     : OFF
Data ex. method : Ethernet
```

---

```
local address: LID 0000 QPN 0x1d26 PSN 0xae3ea9
GID: 00:00:00:00:00:00:00:00:00:00:255:255:192:18:01:02
remote address: LID 0000 QPN 0x2589 PSN 0x554df6
GID: 00:00:00:00:00:00:00:00:00:00:255:255:192:18:01:01
```

---

#bytes	#iterations	BW peak[Gb/sec]	BW average[Gb/sec]	MsgRate[Mpps]
2	1000	0.12	0.12	7.411412
4	1000	0.25	0.23	7.068284
8	1000	0.52	0.46	7.141944
16	1000	1.03	0.91	7.140323
32	1000	1.32	1.31	5.099323
64	1000	3.97	3.62	7.067015
128	1000	8.14	8.12	7.930033
256	1000	15.88	15.24	7.439436
512	1000	31.55	31.40	7.667043
1024	1000	54.00	49.16	6.001562
2048	1000	78.86	73.32	4.475048
4096	1000	91.97	91.61	2.795752
8192	1000	93.04	92.99	1.418913
16384	1000	95.80	95.77	0.730660
32768	1000	96.57	96.57	0.368391
65536	1000	97.04	96.85	0.184723

---

131072	1000	97.25	97.24	0.092740
262144	1000	97.43	97.42	0.046453
524288	1000	97.48	97.48	0.023240
1048576	1000	97.52	97.52	0.011625
2097152	1000	97.69	97.69	0.005823
4194304	1000	97.81	97.80	0.002915
8388608	1000	97.81	97.81	0.001457

---

The previous pair of commands can be repeated for each of the four devices on each DGX-1 host: mlx5\_0, mlx5\_1, mlx5\_2, mlx5\_3. Furthermore, all DGX-1 nodes can be tested pairwise to confirm full RoCE connectivity.

## FlashBlade Configuration

On the **Settings > Network** tab of the FlashBlade UI, or using the corresponding CLI commands, create a subnet for 198.18.0.0/24 as follows:

```
Name:      data-net1
Prefix:    198.18.0.0/24
VLAN:      3000
Gateway:   (leave empty)
MTU:       9000
```

Add an interface for data traffic in 198.18.0.0/24 as follows:

```
Name:      data1
Address:    198.18.0.100
```

From the **Storage** tab, or using the corresponding CLI commands, create a filesystem to use for storing datasets:

```
Name:      datasets
Protocols: enable NFS
```

On each of the DGX-1 servers, mount FlashBlade using the following commands:

```
$ mkdir -p /mnt/datasets
$ mount -t nfs 198.18.0.100:/datasets /mnt/datasets
```

Note that all NFS mount options are left at the OS defaults, which corresponds to a 512kB read- and write-size with FlashBlade and the local file caching option (fsc) disabled.

## ADDITIONAL RESOURCES

- AIRI Github Site: <https://github.com/PureStorage-OpenConnect/AIRI>
- AIRI Reference Architecture: See github site above
- AIRI Product Page: [www.purestorage.com/airi](http://www.purestorage.com/airi)

## APPENDIX A: HOROVOD CONTAINER IMAGE

The following Dockerfile is used to create an image, derived from the base NVIDIA container for Tensorflow, to run Horovod in an OpenMPI environment.

```
FROM nvcr.io/nvidia/tensorflow:17.12

RUN mkdir /build
WORKDIR /build
RUN apt-get update && apt-get install -y --no-install-recommends \
    libibverbs1 \
    libibverbs-dev \
    libmlx5-1 \
    librdmacm-dev \
    librdmacm1 \
    openssh-client \
    openssh-server \
    file \
    && \
    rm -rf /var/lib/apt/lists/*

ENV OPENMPI_VERSION 3.0.0
ENV OPENMPI_TAR openmpi-${OPENMPI_VERSION}.tar.gz
ENV OPENMPI_URL https://www.open-mpi.org/software/ompi/v3.0/downloads

RUN wget -q -O - ${OPENMPI_URL}/${OPENMPI_TAR} | tar -xzf - && \
    cd openmpi-${OPENMPI_VERSION} && \
    ./configure --enable-orterun-prefix-by-default \
        --with-cuda --with-verbs \
        --prefix=/usr/local/mpi --disable-getpwuid && \
    make -j"$(nproc)" install && \
    cd .. && rm -rf openmpi-${OPENMPI_VERSION}
```

```

ENV PATH /usr/local/mpi/bin:$PATH

RUN mkdir -p /var/run/sshd && \
    mkdir -p /root/.ssh && \
    echo "StrictHostKeyChecking no" >> /etc/ssh/ssh_config && \
    echo "UserKnownHostsFile /dev/null" >> /etc/ssh/ssh_config && \
    echo "LogLevel quiet" >> /etc/ssh/ssh_config && \
    sed -i 's/^Port 22/Port 2222/' /etc/ssh/sshd_config && \
    echo "HOST *" > /root/.ssh/config && \
    echo "PORT 2222" > /root/.ssh/config && \
    mkdir -p /root/.ssh && \
    ssh-keygen -t rsa -b 4096 -f /root/.ssh/id_rsa -N "" && \
    cp /root/.ssh/id_rsa.pub /root/.ssh/authorized_keys && \
    chmod 700 /root/.ssh && \
    chmod 600 /root/.ssh/*

RUN export HOROVOD_GPU_ALLREDUCE=NCCL && \
    export HOROVOD_NCCL_INCLUDE=/usr/include && \
    export HOROVOD_NCCL_LIB=/usr/lib/x86_64-linux-gnu && \
    ln -s /usr/local/cuda/lib64/stubs/libcudart.so ./libcudart.so.1 && \
    export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$PWD && \
    pip2 install --no-cache-dir horovod && \
    rm ./libcudart.so.1

```

```
RUN ldconfig
```

```
EXPOSE 2222
```

© 2018 Pure Storage, Inc. All rights reserved.

AIRI, the AIRI logo, Pure Storage, the P Logo, and FlashBlade are trademarks or registered trademarks of Pure Storage, Inc. in the U.S. and other countries. NVIDIA, DGX-1, and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation. All other trademarks are registered marks of their respective owners.

The Pure Storage and NVIDIA products and programs described in this documentation are distributed under a license agreement restricting the use, copying, distribution, and decompilation/reverse engineering of the products. No part of this documentation may be reproduced in any form by any means without prior written authorization from Pure Storage, Inc. and its licensors, if any. Pure Storage and NVIDIA may make improvements and/or changes in the Pure Storage and NVIDIA products and/or the programs described in this documentation at any time without notice.

THIS DOCUMENTATION IS PROVIDED "AS IS" AND ALL EXPRESS OR IMPLIED CONDITIONS, REPRESENTATIONS AND WARRANTIES, INCLUDING ANY IMPLIED WARRANTY OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, OR NON-INFRINGEMENT, ARE DISCLAIMED, EXCEPT TO THE EXTENT THAT SUCH DISCLAIMERS ARE HELD TO BE LEGALLY INVALID. PURE STORAGE SHALL NOT BE LIABLE FOR INCIDENTAL OR CONSEQUENTIAL DAMAGES IN CONNECTION WITH THE FURNISHING, PERFORMANCE, OR USE OF THIS DOCUMENTATION. THE INFORMATION CONTAINED IN THIS DOCUMENTATION IS SUBJECT TO CHANGE WITHOUT NOTICE.

ps\_wp13p\_airi-configuration-guide\_ltr\_01

SALES@PURESTORAGE.COM | 800-379-PURE | @PURESTORAGE

